



Hello!

My name is Seung Woo Choi,
and I am a Data Scientist.

You can find me at:

[linkedin.com/in/seung-woo-choi/](https://www.linkedin.com/in/seung-woo-choi/)



Food Image Classification via CNNs



My Vision





Potential Applications

Recipe Recommender System

The user takes a picture of ANY food with his/her phone, and the model accurately predicts the food class and provides a generic recipe for the food class.

Uses: traveling in a foreign country, trying a new cuisine

Object Classification


The user takes a picture of ANY object with his/her phone, and the model accurately identifies the object.

Use: “eyes” for the visually impaired

Dinner Recommender System

The user takes a picture of his/her dinner, and the model accurately predicts the food class and recommends a similar/different food class for tomorrow’s dinner.

Use: deciding what to eat for dinner tomorrow

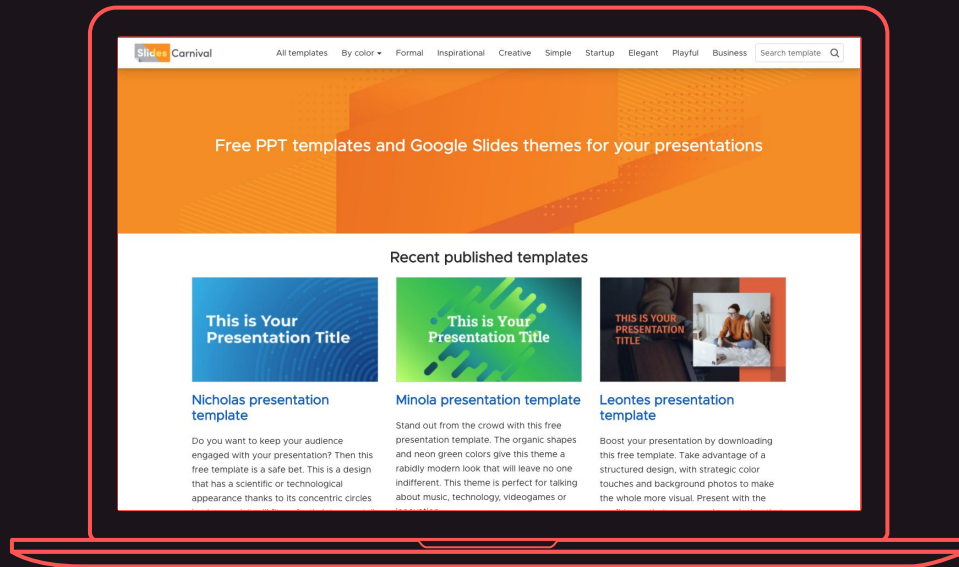


A Reality Check...

Silicon Valley:

Not Hotdog

(Season 4,
Episode 4 Clip)



Problem Statement

- Collect image data for 5 food classes:
 - Burger
 - Hotdog
 - Pizza
 - Taco
 - Sushi
- Build CNNs from scratch and pre-trained, SOTA models to predict multiclass food classes



About the Data

- **Data source:** Reddit's pushshift.io API
- **Collection method:** scraping urls
- **Image data:** .jpg and .png files
- **Ideal sample size:** 1,000 images/class
- **Actual sample size:** 253-282/class
- **Retention rate:** scraped 600/class → used < 300/class → **42-47% retention**
- **Efficiency:** **26.7%** of ideal sample size
- **Data cleaning:** removed manually



Limitations

Data Collection

Google Images vs.
Reddit

Sample Size

Ideal vs. restricted due
to manual removal of
incorrect images

Low Budget

AWS SageMaker

Processing Speed

Test runs for modeling
Hyperparameter
tuning




Exploratory Data Analysis





Class Membership



Example: Burger and Hotdog



Example: Pizza and Taco



Example: Sushi













Image Before and After Resizing

Data Augmentation: Random Rotation



Modeling





Null Model

Majority Class: **Burgers**

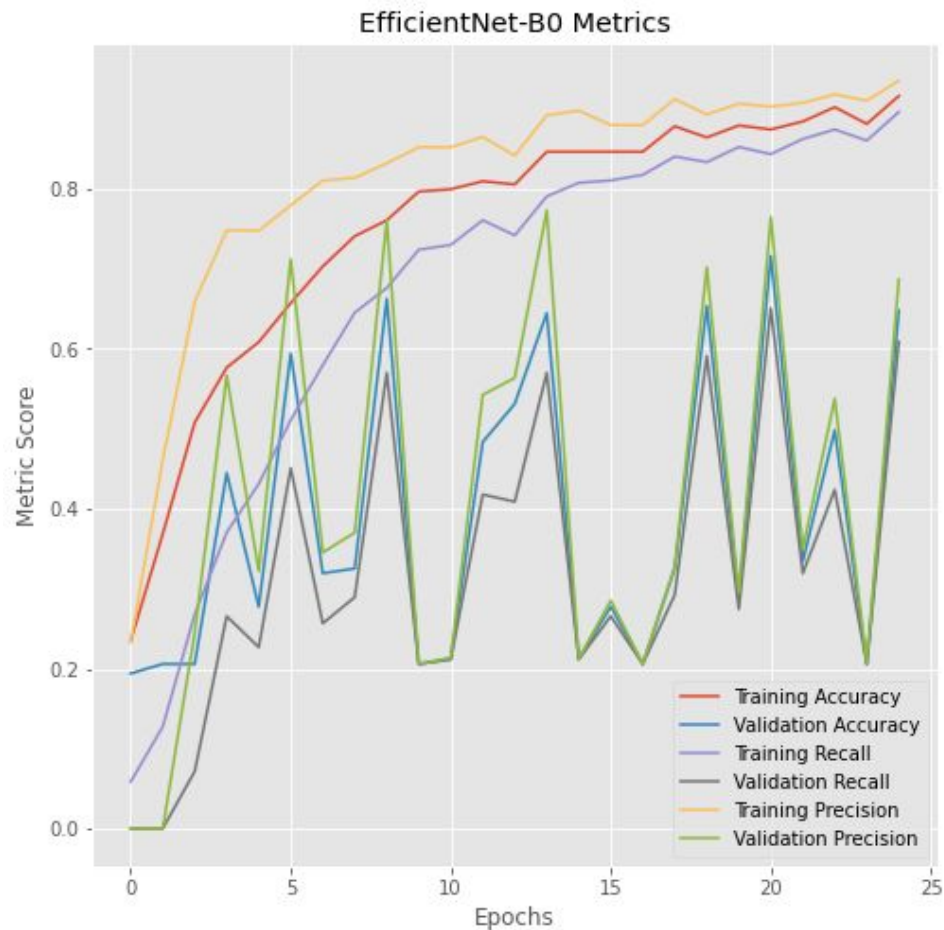
Baseline Accuracy: **21.06%**

Model Type	Validation Accuracy	Vs. Null Model	Validation Loss
CNN with 9 layers	45.37%	2.15x	1.57
CNN with 15 layers	33.13%	1.57x	1.85
ResNet50 (IBM)	28.66%	1.36x	1.57
EfficientNet-B0 (Google)	64.78%	3.07x	1.06

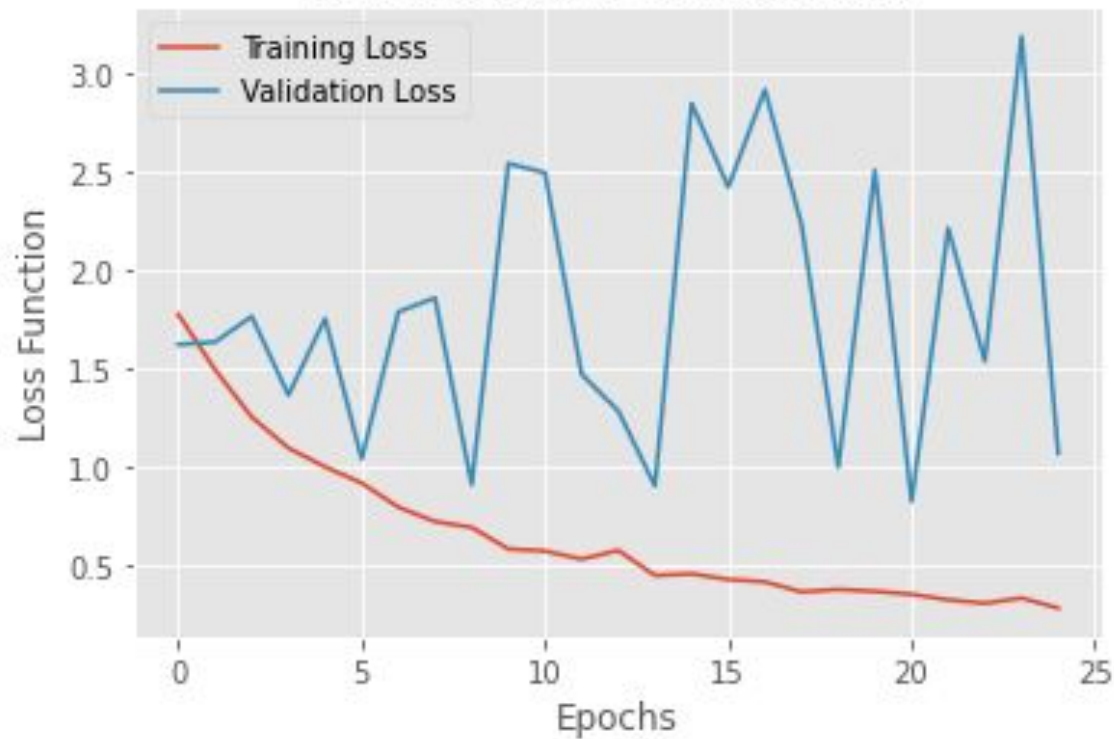
A close-up photograph of four silver spoons arranged diagonally across a dark, textured surface. The spoons contain different types of spices: the first spoon on the left has a reddish-brown powder, the second has a yellowish-brown powder, the third has a dark brown powder, and the fourth has dark, irregularly shaped peppercorns. Scattered around the spoons are various whole spices, including small brown seeds and larger, darker peppercorns. The lighting is dramatic, highlighting the textures of the spices and the metallic sheen of the spoons.

Production Model

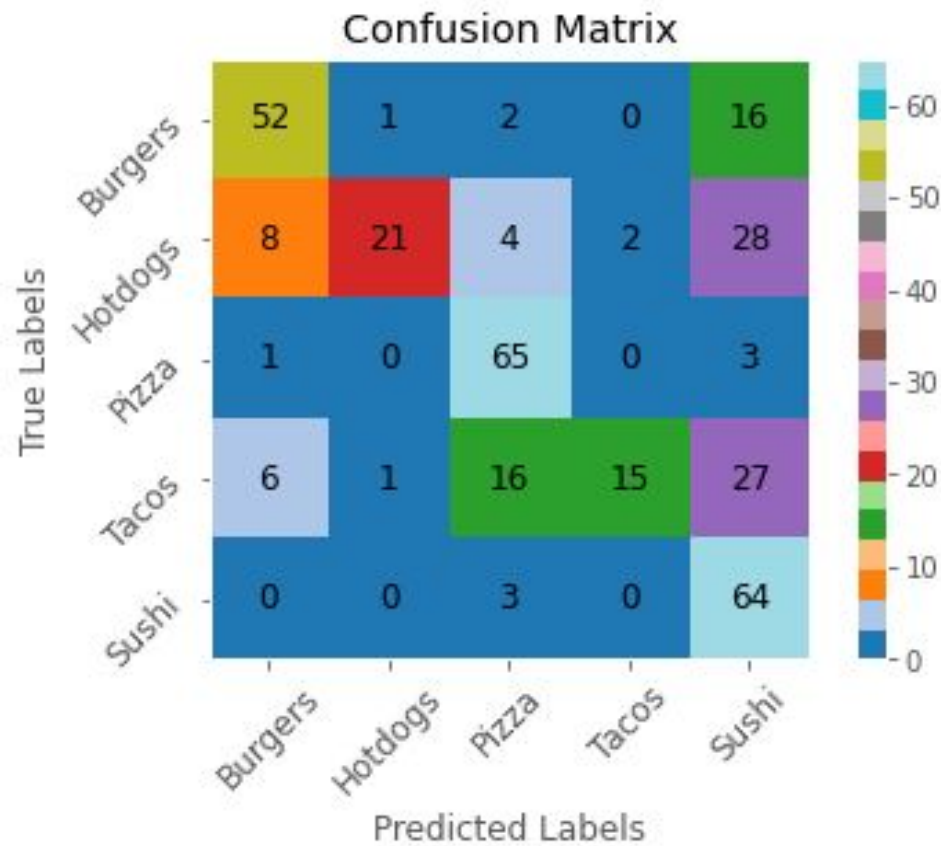
EfficientNet-B0: Classification Metrics



EfficientNet-B0 Loss Function



EfficientNet-B0: Loss Function



EfficientNet-B0: Confusion Matrix

Conclusion



5

Number of classes

253 to 282

Number of images per class

64.78%

Accuracy on unseen data



Thanks!

Any questions?